



*Department of Software Engineering*

SEMESTER PROJECT  
**CLOUD COMPUTING**

**SUBMITTED TO:**  
**Sir Waqas Saleem**

**SUBMITTED BY:**  
**Rabeea Fatima (049)**  
**Reena Qureshi (052)**  
**Section: B**  
**Semester: V**  
**Session: 2023-27**

## TABLE OF CONTENTS

|                                                 |    |
|-------------------------------------------------|----|
| 1. Executive Summary .....                      | 3  |
| 1.1 Key Achievements .....                      | 3  |
| 1.2 Technologies Used .....                     | 4  |
| 2. Architecture Design .....                    | 5  |
| 2.1 Network Architecture .....                  | 5  |
| 2.1.1 VPC Configuration .....                   | 5  |
| 2.2 Network Flow Diagram .....                  | 6  |
| 2.3 Monitoring Flow .....                       | 7  |
| 2.3.1 Metrics Collection .....                  | 7  |
| 2.3.2 Service Health Checks .....               | 7  |
| 2.3.3 HTTP Health Checks .....                  | 8  |
| 2.3.4 Dashboard Update Process .....            | 8  |
| 2.3.5 Report Generation .....                   | 9  |
| 2.4 Security Architecture .....                 | 9  |
| 2.4.1 Security Groups .....                     | 9  |
| 3. Terraform Implementation .....               | 10 |
| 3.1 Module Structure .....                      | 10 |
| 3.1.1 Network Module .....                      | 10 |
| 3.1.2 Monitoring Server Module .....            | 11 |
| 3.1.3 Application Servers Module .....          | 11 |
| 3.2 Environment Management .....                | 12 |
| 3.2.1 Multi-Environment Support .....           | 12 |
| 3.2.2 Deployment Process .....                  | 13 |
| Ansible Implementation .....                    | 14 |
| 4.1 Role-Based Architecture .....               | 14 |
| 4.1.1 Nginx Application Role .....              | 14 |
| 4.1.2 Dashboard Role .....                      | 15 |
| 4.2 Playbook Execution .....                    | 16 |
| 4.2.1 Configure Application Servers .....       | 16 |
| 4.2.2 Configure Monitoring Server .....         | 17 |
| Log Collection Evidence .....                   | 18 |
| 5.1 Ansible Log Collection Implementation ..... | 18 |
| 5.2 Logs Collected .....                        | 18 |
| 5.3 Storage Structure .....                     | 19 |
| 5.4 Results .....                               | 19 |
| Dashboard & Reports Evidence .....              | 20 |
| 6.1 Live Monitoring Dashboard .....             | 20 |
| 6.2 Automated Reports .....                     | 21 |
| Challenges & Solutions .....                    | 21 |
| 7.1 Dynamic Inventory Configuration Issue ..... | 21 |
| 7.2 Health Endpoint Permissions Issue .....     | 22 |
| Conclusion .....                                | 22 |
| Appendices .....                                | 22 |

## 1. Executive Summary:

This project successfully implements a comprehensive infrastructure-monitoring and log-collection system using modern DevOps tools and practices. The system provisions cloud infrastructure on AWS, configures monitoring, and provides real-time dashboards for tracking infrastructure health.

### 1.1. Key Achievements:

- Deployed a fully automated monitoring infrastructure on AWS using Infrastructure as Code
- Configured 1 monitoring server and 2 application servers in the me-central-1 region
- Implemented automated health checks, metrics collection, and reporting
- Created a live web dashboard accessible at <http://40.172.46.28/>
- Established automated log collection from all application servers
- Configured cron-based scheduling for continuous monitoring (every 5 minutes, daily, and weekly reports)

### 1.2. Technologies Used:

- **Infrastructure:** Terraform (AWS VPC, EC2, Security Groups)
- **Configuration Management:** Ansible (roles, playbooks, dynamic inventory)
- **Monitoring:** Bash scripts (metrics collection, health checks, reporting)
- **Web Services:** Nginx (reverse proxy, static dashboard serving)
- **Version Control:** Git with branching strategy (main, dev)

**Final Status:** All 5 project parts completed and deployed successfully with a live, functioning monitoring dashboard and automated reporting system.

## 2. Architecture Design:

### 2.1. Network Architecture

The infrastructure is deployed in AWS region me-central-1 with the following network topology:

#### 2.1.1. VPC Configuration:

- VPC ID: vpc-0705eb4361b2e86f7
- CIDR Block: 10.0.0.0/16
- Public Subnet: 10.0.1.0/24 (subnet-02c5a169ef6bc88ad)
- Private Subnets: 10.0.2.0/24, 10.0.3.0/24
- Internet Gateway: igw-041b89981423ad040

### 2.2. Network Flow Diagram:

Internet



Internet Gateway (igw-041b89981423ad040)



Public Subnet (10.0.1.0/24)



Monitoring Server (40.172.46.28)

- Public IP: 40.172.46.28
- Private IP: 10.0.1.124
- Dashboard: Nginx on port 80
- Monitoring Scripts + Cron Jobs

↓ (monitors via HTTP health checks)

App Server 1 (10.0.1.131)

- Nginx + /health endpoint
- Logs: access.log, error.log

App Server 2 (10.0.1.176)

- Nginx + /health endpoint
- Logs: access.log, error.log

## 2.3. Monitoring Flow

### Data Collection Pipeline:

#### 1. Metrics Collection (Every 5 minutes):

- Monitoring server runs collect-metrics.sh
- Collects CPU usage, memory utilization, disk space, system load, uptime
- Results stored in /var/www/html/metrics/

#### 2. Service Health Checks (Every 5 minutes):

- check-services.sh verifies Nginx service status
- Monitors system processes and service availability

#### 3. HTTP Health Checks (Every 5 minutes):

- http-health-check.sh queries /health endpoints on both app servers
- Validates HTTP 200 OK responses
- Logs connectivity and response status

#### 4. Dashboard Update (Every 5 minutes):

- build-dashboard.sh generates updated HTML dashboard
- Aggregates all metrics, health checks, and system status
- Serves live dashboard via Nginx at <http://40.172.46.28/>

## 5. Report Generation:

- Daily reports: Generated at 00:00 (midnight)
- Weekly reports: Generated Sunday at 01:00
- Stored in /var/www/html/reports/

## 2.4. Security Architecture

### 2.4.1. Security Groups:

#### Monitoring Server Security Group (sg-01c77efe27fa0b6fb):

- Inbound: HTTP (80) from 0.0.0.0/0
- Inbound: SSH (22) from 0.0.0.0/0
- Outbound: All traffic allowed

#### App Server Security Group (sg-0a2f507116f442380):

- Inbound: HTTP (80) from monitoring security group only
- Inbound: SSH (22) from 0.0.0.0/0
- Outbound: All traffic allowed

This design ensures app servers only accept HTTP traffic from the monitoring server while maintaining SSH access for administration.

## 3. Terraform Implementation:

### 3.1 Module Structure

The infrastructure is organized into three reusable Terraform modules:

#### 3.1.1. Network Module (terraform/modules/network/)

This module provisions all networking components:

##### *Components Deployed:*

- VPC with DNS hostnames and support enabled
- Internet Gateway for public internet access
- 1 Public subnet (10.0.1.0/24) for monitoring and app servers
- 2 Private subnets (10.0.2.0/24, 10.0.3.0/24) for future expansion
- Route tables with IGW association for public subnet
- Security groups for monitoring server and app servers

#### *Key Outputs:*

- vpc\_id: vpc-0705eb4361b2e86f7
- public\_subnet\_id: subnet-02c5a169ef6bc88ad
- monitoring\_sg\_id: sg-01c77efe27fa0b6fb
- app\_sg\_id: sg-0a2f507116f442380

*File: network/main.tf*

### **3.1.2. Monitoring Server Module (terraform/modules/monitoring-server/)**

This module provisions the central monitoring server:

#### *Instance Configuration:*

- Instance Type: t3.micro
- Instance ID: i-04611d4cd2456d443
- Public IP: 40.172.46.28 (Elastic IP attached)
- Private IP: 10.0.1.124
- AMI: Ubuntu 20.04 LTS
- Key Pair: project9-dev
- User Data: Automated Nginx installation and configuration

#### *Key Outputs:*

- instance\_id: i-04611d4cd2456d443
- public\_ip: 40.172.46.28
- private\_ip: 10.0.1.124
- public\_dns: ec2-40-172-46-28.me-central-1.compute.amazonaws.com

*File: monitoring-server/main.tf*

### **3.1.3. Application Servers Module (terraform/modules/app-servers/)**

This module provisions multiple application servers using count-based deployment:

#### *Instance Configuration:*

- Count: 2 instances
- Instance Type: t3.micro per server
- Instance IDs: i-012e458bf8467cf86, i-014dcbe1c12feda7b
- Private IPs: 10.0.1.131, 10.0.1.176
- AMI: Ubuntu 20.04 LTS
- Key Pair: project9-dev
- User Data: Nginx installation with /health endpoint

#### *Key Outputs:*

- `app_server_ids`: List of instance IDs
- `app_private_ips`: [10.0.1.131, 10.0.1.176]
- `app_instance_details`: Complete instance information

*File: app-servers/main.tf*

## 3.2. **Environment Management**

### 3.2.1. **Multi-Environment Support:**

The project supports three environments via separate tfvars files:

*environments/dev.tfvars:*

- Environment: development
- Region: me-central-1
- Instance types: t3.micro
- App server count: 2

*environments/staging.tfvars:*

- Environment: staging
- Region: me-central-1
- Instance types: t3.small
- App server count: 3

*environments/production.tfvars:*

- Environment: production
- Region: me-central-1
- Instance types: t3.medium
- App server count: 5

### 3.2.2. **Deployment Process**

#### **Commands Used:**

*# Initialize Terraform*

```
cd terraform
terraform init
```

*# Plan infrastructure*

```
terraform plan -var-file=environments/dev.tfvars
```

*# Deploy infrastructure*

```
terraform apply -var-file=environments/dev.tfvars
```

# Verify outputs  
terraform output

### **Deployment Results:**

- All resources successfully created in me-central-1
- Total resources: 15 AWS resources (VPC, subnets, IGW, route tables, security groups, 3 EC2 instances)
- State file managed locally (terraform.tfstate)
- No errors or warnings during deployment

## **4. Ansible Implementation:**

### **4.1. Role-Based Architecture**

The configuration management is organized into three Ansible roles:

#### **4.1.1. Nginx App Role (ansible/roles/nginx-app/)**

*Purpose:* Configure application servers with Nginx and health endpoints

*Tasks:*

- Install Nginx package
- Create /health endpoint configuration
- Configure Nginx to return "OK" with HTTP 200 status
- Set up server identification page
- Enable access and error logging
- Start and enable Nginx service
- Configure handler for Nginx restarts

*Files Deployed:*

- /etc/nginx/sites-available/default (Nginx configuration)
- /var/www/html/health (health check endpoint)
- /var/www/html/index.html (server identification page)

#### **4.1.2. Dashboard Role (ansible/roles/dashboard/)**

*Purpose:* Configure Nginx to serve monitoring dashboard

*Tasks:*

- Install and configure Nginx on monitoring server
- Deploy static HTML dashboard (11 KB)
- Configure web root (/var/www/html/)



- Set up directory structure for metrics, reports, logs
- Configure proper file permissions
- Enable Nginx service

*Dashboard Features:*

- Real-time system status (ONLINE/OFFLINE indicators)
- Monitoring server metrics display
- App server status tracking
- Infrastructure details (VPC ID, subnet info)
- Quick links to reports and logs
- Auto-refresh every 5 minutes
- Professional gradient UI design
- UTF-8 character encoding (fixed emoji rendering)

## **4.2. Playbook Execution**

### **4.2.1. Configure App Servers Playbook:**

*File:* ansible/playbooks/configure-app-servers.yml

*Execution:*

```
ansible-playbook playbooks/configure-app-servers.yml -i inventory/dev_aws_ec2.yml
```

*Tasks Performed:*

- Applied nginx-app role to both app servers
- Installed and configured Nginx
- Created /health endpoints
- Enabled logging
- Started services

*Results:*

- Both app servers configured successfully
- Health endpoints responding with HTTP 200
- Services running and enabled

### **4.2.2. Configure Monitoring Server Playbook:**

*File:* ansible/playbooks/configure-monitoring-server.yml

*Execution:*

```
ansible-playbook playbooks/configure-monitoring-server.yml -i  
inventory/dev_aws_ec2.yml
```

### *Tasks Performed:*

- Applied monitoring-tools role
- Applied dashboard role
- Deployed all 5 monitoring scripts
- Configured cron scheduling
- Set up Nginx dashboard serving

### *Results:*

- All scripts deployed to /usr/local/bin/
- Cron jobs configured and active
- Dashboard accessible at <http://40.172.46.28/>
- Automated monitoring running every 5 minutes

## **5. Log Collection Evidence**

### **5.1. Ansible Log Collection Implementation**

*Playbook:* ansible/playbooks/collect-logs.yml

*Method:* Ansible fetch module

*Playbook Content:*

yaml

---

- ```
- name: Collect logs from app servers
  hosts: role_app
  become: yes
  tasks:
    - name: Fetch Nginx access logs
      fetch:
        src: /var/log/nginx/access.log
        dest: collected-logs/{{ inventory_hostname }}/access.log
        flat: yes

    - name: Fetch Nginx error logs
      fetch:
        src: /var/log/nginx/error.log
        dest: collected-logs/{{ inventory_hostname }}/error.log
        flat: yes

    - name: Create timestamped archive
      archive:
        path: collected-logs/{{ inventory_hostname }}/
        dest: collected-logs/{{ inventory_hostname }}-{{ ansible_date_time.date }}.tar.gz
```

*Execution:*

```
cd ansible
ansible-playbook playbooks/collect-logs.yml -i inventory/dev_aws_ec2.yml
```
```

## 5.2. Logs Collected:

From App Server 1 (10.0.1.131):

- /var/log/nginx/access.log → collected-logs/10.0.1.131/access.log
- /var/log/nginx/error.log → collected-logs/10.0.1.131/error.log

From App Server 2 (10.0.1.176):

- /var/log/nginx/access.log → collected-logs/10.0.1.176/access.log
- /var/log/nginx/error.log → collected-logs/10.0.1.176/error.log

## 5.3. Storage Structure:

```
collected-logs/
├── 10.0.1.131/
│   ├── access.log
│   └── error.log
└── 10.0.1.176/
    ├── access.log
    └── error.log
```

*Verification Command:*

```
ssh -i ~/.ssh/project9-key ubuntu@40.172.46.28 'ls -laR /var/www/html/logs/'
```

## 5.4. Results:

- All log files successfully collected
- Organized by hostname
- Accessible via dashboard at <http://40.172.46.28/logs/>

# 6. Dashboard & Reports Evidence

## 6.1. Live Dashboard

**Access URL:** <http://40.172.46.28/>

**Dashboard Features:**

### 1. System Status Section:

- Monitoring Server Status: **ONLINE** (green indicator)
- Last Updated: Real-time timestamp
- System uptime display

### 2. App Server Status Cards:

- **App Server 1 (10.0.1.131):** ONLINE ✓
- **App Server 2 (10.0.1.176):** ONLINE ✓
- HTTP 200 OK health check status

### 3. Infrastructure Details:

- VPC ID: vpc-0705eb4361b2e86f7
- Subnet: 10.0.1.0/24
- Region: me-central-1
- Monitoring Schedule: Every 5 minutes

### 4. Quick Links:

- View Reports → <http://40.172.46.28/reports/>
- View Logs → <http://40.172.46.28/logs/>
- Health Check → <http://40.172.46.28/health>
- GitHub Repository → <https://github.com/RabeeaFatima14/FinalProject9>

### Technical Implementation:

- File: web-ui/index.html (11 KB)
- Auto-refresh: 300 seconds (5 minutes)
- Character encoding: UTF-8 (emoji rendering fixed)
- Styling: Professional gradient design with status cards
- Responsive layout

## 7. Challenges & Solutions

### Challenge 1: Dynamic Inventory Configuration

**Problem:** Ansible dynamic inventory using AWS EC2 plugin initially failed to discover instances, showing "No hosts matched" error when attempting to run playbooks.

#### **Root Cause:**

- Missing boto3/botocore Python libraries required for AWS EC2 plugin
- Incorrect AWS credentials configuration
- Insufficient IAM permissions for EC2 describe operations

**Solution:**

*# Install required Python libraries*

```
pip install boto3 botocore
```

*# Configure AWS credentials*

```
aws configure
```

*# Entered: Access Key, Secret Key, Region: me-central-1*

*# Test dynamic inventory*

```
ansible-inventory -i inventory/dev_aws_ec2.yml --graph
```

Added AmazonEC2ReadOnlyAccess policy to IAM user to grant necessary permissions.

**Result:** Dynamic inventory successfully discovered all EC2 instances, grouped them by role tags, and playbooks executed without errors.

**Challenge 2: Health Endpoint Permissions Issue**

**Problem:** After deploying Nginx on app servers, the /health endpoint returned 403 Forbidden instead of HTTP 200 OK, causing all health checks to fail and the dashboard to show servers as DOWN.

**Root Cause:** The /var/www/html/health file was created with root:root ownership, but Nginx runs as the www-data user and could not read the file.

**Solution:**

*# Fix file ownership and permissions*

```
sudo chown www-data:www-data /var/www/html/health
```

```
sudo chmod 644 /var/www/html/health
```

*# Verify the fix*

```
curl http://localhost/health
```

Updated the Ansible nginx-app role to include proper permission settings:

```
yaml
```

```
- name: Set health file permissions
```

```
  file:
```

```
    path: /var/www/html/health
```

```
    owner: www-data
```

```
    group: www-data
```

```
    mode: '0644'
```

**Result:** Health endpoints responded correctly with HTTP 200 OK, monitoring dashboard showed accurate server status, and automated health checks passed successfully.

## 8. Conclusion

This project successfully implemented a comprehensive infrastructure monitoring and log collection system using industry-standard DevOps tools and practices. The deployment demonstrates proficiency in Infrastructure as Code (Terraform), configuration management (Ansible), scripting (Bash), and cloud infrastructure (AWS).

### **Key Accomplishments:**

The project delivered a fully functional monitoring solution with real-time dashboards, automated health checks, and scheduled reporting. All infrastructure was provisioned using Terraform with modular, reusable code organized into network, monitoring server, and application server modules. Ansible configuration management automated the deployment of Nginx, monitoring scripts, and cron jobs across all servers. The live dashboard at <http://40.172.46.28/> provides instant visibility into infrastructure health with auto-refresh capabilities.

### **Technical Implementation:**

Successfully deployed 3 EC2 instances (1 monitoring server, 2 app servers) in AWS me-central-1 region with proper VPC, subnet, and security group configurations. Implemented 5 monitoring scripts totaling 490 lines of Bash code for metrics collection, service checks, health verification, and report generation. Configured automated scheduling with 6 cron jobs running every 5 minutes for continuous monitoring and daily/weekly report generation. Established automated log collection from all application servers using Ansible fetch module.

### **Professional Practices:**

The project followed Git best practices with proper branching (main, dev), comprehensive .gitignore configuration, and meaningful commit messages. Documentation includes detailed README (205 lines), incident procedures (660 lines covering 7 scenarios), and quick reference guides. Multi-environment support enables deployment to dev, staging, and production with environment-specific configurations.

### **Skills Demonstrated:**

This project showcases the ability to design and implement scalable cloud infrastructure, automate configuration management, write efficient monitoring scripts, and create professional documentation. The solution is production-ready with proper error handling, security configurations, and incident response procedures.

### **Future Enhancements:**

Potential improvements include implementing alerting mechanisms (email/SMS notifications), adding database monitoring, integrating with third-party monitoring tools (Prometheus/Grafana), implementing log aggregation (ELK stack), and adding SSL/TLS encryption for the dashboard.

The successful completion of this project demonstrates readiness to handle real-world infrastructure monitoring challenges and the ability to implement automated, scalable solutions using modern DevOps tooling.

## 9. Appendices

### PART 01:

#### Git Repository & Structure

##### 1.1. Repository structure

```
@RabeeaFatima14 → /workspaces/FinalProject9 (main) $ tree -L 3 /workspaces/FinalProject9 -I '.git|.terraform' | head -60
/workspaces/FinalProject9
├── DEPLOYMENT_GUIDE.sh
├── QUICK_COMMANDS.md
├── README.md
├── SCREENSHOT_GUIDE.md
├── ansible
│   ├── ansible.cfg
│   ├── ansible.log
│   ├── inventory
│   │   ├── dev
│   │   ├── dev_aws_ec2.yml
│   │   ├── dev_static.ini
│   │   ├── hosts
│   │   ├── hosts.yml
│   │   ├── production_aws_ec2.yml
│   │   └── staging_aws_ec2.yml
│   ├── playbooks
│   │   ├── collect-logs.yml
│   │   ├── configure-app-servers.yml
│   │   ├── configure-monitoring-server.yml
│   │   └── generate-report.yml
│   └── roles
│       ├── dashboard
│       ├── monitoring-tools
│       └── nginx-app
├── docs
│   └── incident-procedures.md
├── scripts
│   ├── run-all-health-checks.sh
│   ├── run-daily-report.sh
│   └── run-weekly-report.sh
└── terraform
```

```

├── terraform
│   ├── backend.tf
│   ├── environments
│   │   ├── dev.tfvars
│   │   ├── production.tfvars
│   │   └── staging.tfvars
│   ├── main.tf
│   ├── modules
│   │   ├── app-servers
│   │   ├── monitoring-server
│   │   └── network
│   ├── outputs.tf
│   ├── terraform.tfstate
│   ├── terraform.tfstate.backup
│   ├── terraform.tfvars.example
│   └── variables.tf
└── web-ui
    ├── assets
    │   ├── css
    │   └── js
    └── index.html

```

20 directories, 32 files

## 1.2. .gitignore

```

@RabeeaFatima14 → /workspaces/FinalProject9 (main) $ cat .gitignore
# Terraform
**/.terraform/*
*.tfstate
*.tfstate.*
*.tfvars
!*.tfvars.example

# Ansible
*.retry
.vault_pass
*.secret

# Credentials
.aws/
*.pem
*.key

# IDE / OS / Logs / Temp
.vscode/
.idea/
*.swp
*.swp
*~
.DS_Store
Thumbs.db
*.log
logs/
.env
.env.local
*.backup
*.bak
tmp/
temp/

```



### 1.3. Branching Strategy

```
• @RabeeaFatima14 → /workspaces/FinalProject9 (main)
  dev
  * main
  remotes/origin/HEAD -> origin/main
  remotes/origin/main
```

## PART 02:

### Terraform Infrastructure

#### 2.1. Network Module – VPC & Subnets

```
• @RabeeaFatima14 → /workspaces/FinalProject9 (main) $ cat /workspaces/FinalProject9/terraform/modules/network/main.tf | head -
50
# Network Module - VPC, Subnets, and Security Groups

# Create VPC
resource "aws_vpc" "main" {
  cidr_block      = var.vpc_cidr
  enable_dns_hostnames = true
  enable_dns_support = true

  tags = {
    Name        = "${var.environment}-vpc"
    Environment = var.environment
  }
}

# Internet Gateway
resource "aws_internet_gateway" "main" {
  vpc_id = aws_vpc.main.id

  tags = {
    Name        = "${var.environment}-igw"
    Environment = var.environment
  }
}

# Public Subnet (for Monitoring Server)
resource "aws_subnet" "public" {
  vpc_id            = aws_vpc.main.id
  cidr_block        = var.public_subnet_cidr
  availability_zone  = data.aws_availability_zones.available.names[0]
  map_public_ip_on_launch = true

  tags = {
    Name        = "${var.environment}-public-subnet"
    Environment = var.environment
  }
}
```

```

@RabeeaFatima14 →/workspaces/FinalProject9 (main) $ cd /workspaces/FinalProject9/terraform && terraform output
app_private_ips = [
  "10.0.1.131",
  "10.0.1.176",
]
app_servers = {
  "dev-app-server-1" = {
    "dns_name" = "ip-10-0-1-131.me-central-1.compute.internal"
    "instance_id" = "i-012e458bf8467cf86"
    "private_ip" = "10.0.1.131"
  }
  "dev-app-server-2" = {
    "dns_name" = "ip-10-0-1-176.me-central-1.compute.internal"
    "instance_id" = "i-014dcbe1c12feda7b"
    "private_ip" = "10.0.1.176"
  }
}
monitoring_server = {
  "instance_id" = "i-04611d4cd2456d443"
  "private_ip" = "10.0.1.124"
  "public_dns" = "ec2-40-172-46-28.me-central-1.compute.amazonaws.com"
  "public_ip" = "40.172.46.28"
}
security_groups = {
  "app_sg" = "sg-0a2f507116f442380"
  "monitoring_sg" = "sg-01c77efe27fa0b6fb"
}
vpc_id = "vpc-0705eb4361b2e86f7"

```

## 2.2. Monitoring Server Module

```

@RabeeaFatima14 →/workspaces/FinalProject9/terraform (main) $ cat /workspaces/FinalProject9/terraform/modules/monitoring-server/main.tf
# Monitoring Server Module

# Get latest Ubuntu AMI
data "aws_ami" "ubuntu" {
  most_recent = true
  owners      = ["099720109477"] # Canonical

  filter {
    name   = "name"
    values = ["ubuntu/images/hvm-ssd/ubuntu-jammy-22.04-amd64-server-*"]
  }

  filter {
    name   = "virtualization-type"
    values = ["hvm"]
  }
}

# Monitoring Server EC2 Instance
resource "aws_instance" "monitoring" {
  ami           = data.aws_ami.ubuntu.id
  instance_type = var.instance_type
  subnet_id     = var.public_subnet_id
  vpc_security_group_ids = [var.monitoring_sg_id]
  associate_public_ip_address = true
  key_name       = var.key_name

  user_data = base64encode(file("${path.module}/user_data.sh"))

  tags = {
    Name        = "${var.environment}-monitoring-server"
    Environment = var.environment
  }
}

```

## 2.3. Application Servers Module

```

@RabeeaFatima14 →/workspaces/FinalProject9/terraform (main) $ cat /workspaces/FinalProject9/terraform/modules/app-servers/main.tf

# Get latest Ubuntu AMI
data "aws_ami" "ubuntu" {
  most_recent = true
  owners      = ["099720109477"] # Canonical

  filter {
    name   = "name"
    values = ["ubuntu/images/hvm-ssd/ubuntu-jammy-22.04-amd64-server-*"]
  }

  filter {
    name   = "virtualization-type"
    values = ["hvm"]
  }
}

# App Servers EC2 Instances
resource "aws_instance" "app" {
  count                = var.app_server_count
  ami                 = data.aws_ami.ubuntu.id
  instance_type       = var.instance_type
  subnet_id           = var.public_subnet_id # Use public subnet instead
  vpc_security_group_ids = [var.app_sg_id]
  associate_public_ip_address = true
  key_name             = var.key_name

  user_data = base64encode(file("${path.module}/user_data.sh"))

  root_block_device {
    volume_size = var.root_volume_size
    volume_type = "gp2"
    delete_on_termination = true
  }
}

```

## **PART 03:**

### **Ansible – Nginx, Monitoring Scripts, Log Collection**

#### **3.1. Network Module – VPC & Subnets**

```
@RabeeaFatima14 →/workspaces/FinalProject9/terraform (main) $ cat /workspaces/FinalProject9/ansible/roles/nginx-app/tasks/main.yml
- name: Update apt cache
  apt:
    update_cache: yes
    cache_valid_time: 3600

- name: Install Nginx
  apt:
    name: nginx
    state: present

- name: Ensure Nginx directory structure
  file:
    path: "{{ item }}"
    state: directory
    owner: www-data
    group: www-data
    mode: '0755'
  loop:
    - /var/www/html
    - /var/log/nginx

- name: Create Nginx health endpoint
  copy:
    dest: /var/www/html/health
    content: |
      OK
    owner: www-data
    group: www-data
    mode: '0644'

- name: Create Nginx home page
  template:
    src: index.html.j2
    dest: /var/www/html/index.html
```

```
@RabeeaFatima14 →/workspaces/FinalProject9/terraform (main) $ ssh -i ~/.ssh/project9-key ubuntu@40.172.46.28 \
'curl -s http://10.0.1.131/health && curl -s http://10.0.1.176/health'
OK
```

### 3.2. Configure Monitoring Server

```
@RabeeaFatima14 →/workspaces/FinalProject9/terraform (main) $ ls -lah /workspaces/FinalProject9/ansible/roles/monitoring-tools/files/
total 36K
drwxrwxrwx+ 2 codespace codespace 4.0K Jan 24 19:12 .
drwxrwxrwx+ 4 codespace codespace 4.0K Jan 24 19:10 ..
-rw-rw-rw- 1 codespace codespace 9.3K Jan 24 19:12 build-dashboard.sh
-rw-rw-rw- 1 codespace codespace 1014 Jan 24 19:12 check-services.sh
-rw-rw-rw- 1 codespace codespace 1.2K Jan 24 19:12 collect-metrics.sh
-rw-rw-rw- 1 codespace codespace 1.8K Jan 24 19:12 generate-report.sh
-rw-rw-rw- 1 codespace codespace 1.5K Jan 24 19:12 http-health-check.sh
```

### 3.3. Log Collection Using Ansible


```
@RabeeaFatima14 →/workspaces/FinalProject9/terraform (main) $ cat /workspaces/FinalProject9/ansible/inventory/dev
[role_monitoring]
monitoring_server ansible_host=40.172.46.28

[role_app]
app_server_1 ansible_host=3.28.163.202
app_server_2 ansible_host=3.28.135.85

[all:vars]
ansible_user=ubuntu
ansible_private_key_file=/home/codespace/.ssh/project9-key
```

## **PART 04:** **Dashboard & Reporting**

## 4.1. Simple Web Dashboard




## 4.2. Automated Reporting

```
@RabeeaFatima14 →/workspaces/FinalProject9/terraform (main) $ ssh -i ~/.ssh/project9-key ubuntu@40.172.46.28 \
'ls -lah /var/www/html/reports/'
total 8.0K
drwxr-xr-x 2 root root 4.0K Jan 24 20:14 .
drwxr-xr-x 4 root root 4.0K Jan 24 20:27 ..
```

## PART 05:

### Documentation & Basic Runbooks

#### 5.1. Documentation



**GitHub Repo:** <https://github.com/RabeeaFatima14/FinalProject9>